

# Advanced Machine Learning: Data Science and ML Refresher

Shubhankar Agrawal

## Abstract

This document serves as a quick refresher for Data Science and Machine Learning interviews. It covers mathematical and technical concepts across a range of algorithms. This requires the reader to have a foundational level knowledge with tertiary education in the field. This PDF contains material for revision over key concepts that are tested in interviews.

## Contents

<b>1</b>	<b>Deep Learning</b>	<b>1</b>
1.1	Key Concepts . . . . .	1
	Terminology • Regularization	
1.2	More Components . . . . .	1
	Activations • Optimizers • Schedulers • Code	
1.3	Convolutional Neural Networks . . . . .	2
1.4	Recurrent Neural Networks . . . . .	2
	LSTM • GRU • Seq 2 Seq	
1.5	Transformers . . . . .	3
	Attention	
1.6	Unsupervised Approaches . . . . .	3
	Variational Auto Encoder	
<b>2</b>	<b>Natural Language Processing</b>	<b>3</b>
2.1	Terminology . . . . .	3
2.2	Preprocessing . . . . .	3
2.3	Embedding . . . . .	3
	Vectorization • Word Level • Sentence Level	
2.4	Feature Retrieval . . . . .	4
	Named Entity Recognition • POS Tagging	
2.5	Other Models . . . . .	4
	Natural Language Inference • Topic Modelling	
<b>3</b>	<b>Time Series</b>	<b>5</b>
3.1	Analyses . . . . .	5
	Decomposition • Stationarity • Autocorrelation	
3.2	Forecasting . . . . .	5
	Exponential Smoothing • ARIMA • GARCH	

## 1. Deep Learning

### 1.1. Key Concepts

Neural Networks - Architecture similar to neurons




Figure 1. Neural Networks  
[2]

#### 1.1.1. Terminology

Commonly used terms

**Model:** Function  $y = f(x)$  parametrized by  $\theta$

**Epoch:** One full pass of the training set

**Batch:** Subset of the training set

**Back-propagation:** Weigh updates with gradients

Table 1. Gradient Descent

Name	Samples	Pros	Cons
Batch	All	Stable Convergence	Memory, Slow
Mini-Batch	Subset	Efficient, Parallel	Tune Size
Stochastic	1	Escape local minima	Noisy

Layer sizes commonly used  $2^n$ , optimized for GPU computations and memory allocations

#### 1.1.2. Regularization

Common methods:

- L1 - Lasso
- L2 - Ridge
- Data Augmentation (Scale, Rotate, Flip, Noise)

#### Dropout

- Train: Randomly set neurons to 0 (probability  $p$ )
- Infer: Scale activations by  $1 - p$

**Early Stopping:** Stop when validation loss plateaus

**Batch Normalization:** Stabilize training, convergence

$$\hat{x}_i = \frac{x_i - \mu}{\sqrt{\sigma^2 + \epsilon}} \quad (1)$$

## 1.2. More Components

### 1.2.1. Activations

Introduce Non-Linearity, summarized in Figure 5.

Notes:

- ReLU / Leaky ReLU usually used in between layers
- ReLU can cause dead neurons (Solve with Leaky)
- Tanh used to centre outputs around 0
- Sigmoid for Binary output
- Softmax for multiclass output

**Vanishing Gradient:** Activation derivative tends to 0. Causes neurons to turn off. Switch to other activations (ReLU).

**Exploding Gradient:** Gradient becomes too large while propagating. De-stabilizes convergence. Use gradient clipping, Batch Normalization.

**Gradient Checking:** Check analytical (Back-propagated) vs Approximate gradient. The smaller the value the better

$$\nabla_{\theta} J(\theta) \approx \frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2\epsilon} \quad (2)$$
$$\text{Difference} = \frac{\|\nabla_{\theta} J_{\text{analytical}} - \nabla_{\theta} J_{\text{approximate}}\|}{\|\nabla_{\theta} J_{\text{analytical}}\| + \|\nabla_{\theta} J_{\text{approximate}}\|}$$

**NOTE:** Batch Normalization should be applied after ReLU otherwise the normalization is lost after the activation

### 1.2.2. Optimizers

#### Gradient Descent

$$\theta_t = \theta_{t-1} - \eta \nabla_{\theta} J(\theta_{t-1}) \quad (3)$$

#### Momentum

Exponential decay on moving average of gradients

$$\begin{aligned} v_t &= \beta v_{t-1} + (1 - \beta) \nabla_{\theta} J(\theta_{t-1}) \\ \theta_t &= \theta_{t-1} - \eta v_t \end{aligned} \quad (4)$$

#### AdaGrad

Scale by inverse square root of running average of squared gradients.

$$\begin{aligned} s_t &= s_{t-1} + (\nabla_{\theta} J(\theta_{t-1}))^2 \\ \theta_t &= \theta_{t-1} - \frac{\eta}{\sqrt{s_t + \epsilon}} \nabla_{\theta} J(\theta_{t-1}) \end{aligned} \quad (5)$$

#### RMSProp

Improves Adagrad with a Decay factor to prevent fast diminishing weights

$$\begin{aligned} s_t &= \beta s_{t-1} + (1 - \beta) (\nabla_{\theta} J(\theta_{t-1}))^2 \\ \theta_t &= \theta_{t-1} - \frac{\eta}{\sqrt{s_t + \epsilon}} \nabla_{\theta} J(\theta_{t-1}) \end{aligned} \quad (6)$$

#### Adam (Adaptive Moment Estimation)

Combines Momentum and RMSProp

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) \nabla_{\theta} J(\theta_{t-1}) \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) (\nabla_{\theta} J(\theta_{t-1}))^2 \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t} \\ \theta_t &= \theta_{t-1} - \frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t \end{aligned} \quad (7)$$

### 1.2.3. Schedulers

Adapt learning rate over epochs for faster convergence and optimized search

**Table 2.** Learning Rate Schedulers

Method	Description
<b>StepLR</b>	Reduce by constant factor $n$ iterations
<b>ExponentialLR</b>	Reduce exponentially
<b>CyclicLR</b>	Cycle from base to minimum
<b>ReduceLROnPlateau</b>	Reduce when the metric plateaus

### 1.2.4. Code

```
model.train(True)
optimizer.zero_grad()
outputs = model(inputs)
loss = loss_fn(outputs, labels)

loss.backward()
optimizer.step()

model.eval()
with torch.no_grad():
    outputs = model(inputs)
```

**Code 1.** PyTorch Model

### 1.3. Convolutional Neural Networks

Modelling spatial features.

Input Dimensions: ( $BatchSize \times HeightPixels \times WidthPixels$ )

$$(I * K)(x, y) = \sum_m \sum_n I(x + m, y + n) \cdot K(m, n) \quad (8)$$

**Table 3.** CNN Layers

Layer	Description	Purpose
Convolutional	Kernel Multiplication	Spatial Features
Pooling	Aggregates (Max / Avg)	Down Sample
Fully Connected	Weight Multiplication	Learning

$$\text{Output Size} = \frac{(W - K + 2P)}{S} + 1 \quad (9)$$

where:

- $W$  is the input size (width or height),
- $K$  is the kernel size,
- $P$  is the padding applied,
- $S$  is the stride of the convolution.

#### ResNet

CNN with skip connections (Concat identity + convolution features). Solves vanishing gradient.

### 1.4. Recurrent Neural Networks

Modelling sequential (or) temporal features.

Input Dimensions: ( $BatchSize \times Features \times TimeSteps$ )

- Maintains a hidden state  $h_t$  to store temporal influence

$$h_t = f(W_h \cdot h_{t-1} + W_x \cdot x_t + b) \quad (10)$$

Batch Learning

- Process encoding in parallel
- Hidden state updates sequentially

Truncated Back Propagation Through Time (BPTT): Limit propagating gradients through full sequence. Prevent Vanishing Gradients.

**Inference:** No teacher forcing, i.e. use  $\hat{y}_t$  as inputs  $x_{t+1}$

#### 1.4.1. LSTM

Long Short Term Memory

- Two hidden states (Short Term, Long Term)
- Three gates (Input, Forget, Output)

$$\text{Forget } f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

$$\text{Input } i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\text{Cell (Long)} C_t = f_t \cdot C_{t-1} + i_t \cdot \tanh(W_c \cdot [h_{t-1}, x_t] + b_c) \quad (11)$$

$$\text{Output } o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$\text{Hidden (Short)} h_t = o_t \cdot \tanh(C_t)$$

#### 1.4.2. GRU

Gated Recurrent Unit: Simplified version of LSTM with fewer parameters

- Single hidden state
- Two gates (Update and Reset)

$$\begin{aligned}
 \text{Update } z_t &= \sigma(W_z \cdot [h_{t-1}, x_t] + b_z) \\
 \text{Reset } r_t &= \sigma(W_r \cdot [h_{t-1}, x_t] + b_r) \\
 \text{Hidden } h_t &= z_t \cdot h_{t-1} \\
 &\quad + (1 - z_t) \cdot \tanh(W_h \cdot [r_t \cdot h_{t-1}, x_t] + b_h)
 \end{aligned} \tag{12}$$

RNN architectures summarized in Figure 6

#### 1.4.3. Seq 2 Seq

Sequence to sequence models: Input Sequence -> Output Sequence

### 1.5. Transformers

Seq 2 Seq models with an **Encoder** and **Decoder** architecture.

- Replace RNNs with self-attention
- Process sequences in parallel

#### 1.5.1. Attention

Identifies relevant parts of sequence for each token

**Scaled Dot-Product Attention:**

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \tag{13}$$

where  $Q$ : Query,  $K$ : Key,  $V$ : Value,  $d_k$ : Dimensions

**Multi-Head Attention:**

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \tag{14}$$

- Computes attention over multiple independent heads
- Each head learns unique representations

**Types of Attention**

- Self-Attention: Within the same sequence
- Cross-Attention: Between encoder and decoder sequences
- Masked Self-Attention: Used in autoregressive tasks (e.g., GPT) to prevent looking at future tokens

**Input Representations**

- **Token Embeddings:** Maps tokens (words, subwords) to vectors
- **Positional Encodings:** Adds order information to tokens

$$\begin{aligned}
 PE_{\text{pos}, 2i} &= \sin\left(\frac{\text{pos}}{10000^{2i/d_{\text{model}}}}\right), \\
 PE_{\text{pos}, 2i+1} &= \cos\left(\frac{\text{pos}}{10000^{2i/d_{\text{model}}}}\right)
 \end{aligned} \tag{15}$$

**Layers**

- **Encoder:**
  - Layer Norm, Multi-Head Attention, Feedforward Network.
  - Produces sequence representations for input.
- **Decoder:**
  - Masked Self-Attention, Cross-Attention, Feedforward Network.
  - Generates output sequence step by step.


The architecture is demonstrated in Figure 7.

### 1.6. Unsupervised Approaches

#### 1.6.1. Variational Auto Encoder

Generative model to identify underlying distribution

Learns distribution parameters with sampled errors through an encoder decoder framework



**Figure 2.** Variational Auto Encoder [6]

- Generate lower dimensional embeddings
- Remove biases from data
- Identify outliers by checking encodings

## 2. Natural Language Processing

### 2.1. Terminology

**N-grams:** Phrases of n tokens

**Dictionary:** Set of all words

**Document:** Single sample of text

**Corpus:** Set of all documents

**Stop Words:** commonly used filler words (removed)

### 2.2. Preprocessing

Common steps to prepare text

**Tokenization:** Split into words, sub words

**Stemming:** Short form (improving -> improv)

**Lemmatization:** Root word (improving -> improve)

### 2.3. Embedding

#### 2.3.1. Vectorization

Capture occurrences and frequency of appearances

**Bag of Words:** Frequency counts of words

**One Hot Encoding:** Binary Vectors for Tokens

**TF-IDF:** Term Frequency · Inverse Document Frequency

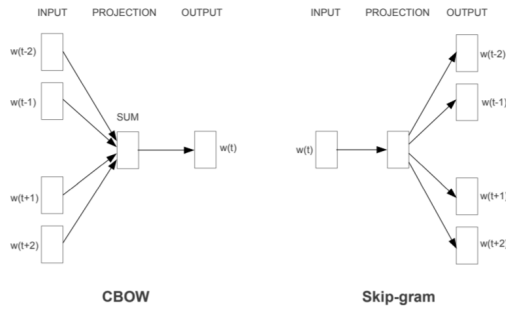
#### TF-IDF

$$\begin{aligned}
 \text{TF}(t, d) &= \frac{\text{Number of times term } t \text{ appears in document } d}{\text{Total number of terms in document } d} \\
 \text{IDF}(t) &= \log\left(\frac{N}{|\{d \in D : t \in d\}|}\right) \\
 \text{TF-IDF}(t, d) &= \text{TF}(t, d) \cdot \text{IDF}(t)
 \end{aligned} \tag{16}$$

#### 2.3.2. Word Level

Capture semantic representations

**Word2Vec:** ML Model for Embeddings



**Figure 3.** Word2Vec Models  
[7]

**Table 4.** Word2Vec

	<b>CBOW</b>	<b>Skip Gram</b>
Input	Context	Token
Output	Token	Context
Embeddings	Input Layer	Output Layer
Pros	Faster	Smaller Datasets
Cons	Rare words	Slower

Negative Sampling: Add a few unrelated target words to reduce overall updates. Model better to distinguish.

**GLoVe**: Global Vectors for Word Representation.  
Factorization of co-occurrence matrix

- Preserves global corpus
- Computationally efficient

### 2.3.3. Sentence Level

**BERT**: Bidirectional Encoder Representations from Transformers  
Relies on Bidirectional Attention

- Masked Language Modelling (MLM)  
Randomly mask tokens and predict probability

$$\mathcal{L}_{\text{MLM}} = - \sum_{i \in M} \log P(x_i | \hat{x}) \quad (17)$$

- Next Sentence Prediction (NSP)  
Predict if a sentence follows another

$$\mathcal{L}_{\text{NSP}} = - \sum_i y_i \log P(y_i) + (1 - y_i) \log(1 - P(y_i)) \quad (18)$$

[CLS] token at the beginning of the sentence is a pooled representation.

#### Sentence BERT

- Generates Fixed Size Sentence embeddings
- Relies on Siamese Networks

Siamese Networks

- Two Sentences passed through the same BERT model
- Similarity over pooled sentence embeddings

Contrastive Loss

- Pairs of Sentences
- Calculate distance ( $d$ ) between embeddings
- Classify labels ( $y$ ) [1 Similar, -1 dissimilar]

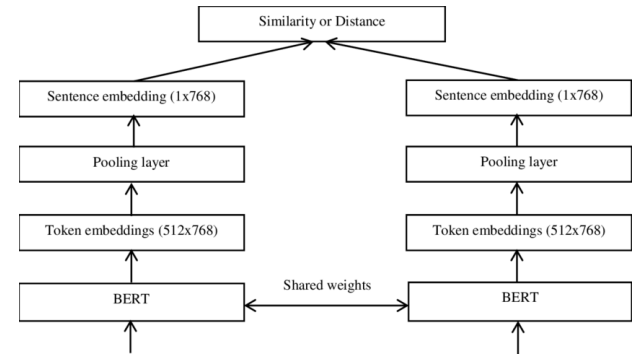
$$\mathcal{L}_{\text{contrastive}} = \frac{1}{2} (y \cdot d^2 + (1 - y) \cdot \max(0, \text{margin} - d)^2) \quad (19)$$

Triplet Loss

- Anchor ( $u$ ), with positive ( $v^+$ ) and negative ( $v^-$ ) sentences
- Similarity scores (Cosine distance)

$$\mathcal{L}_{\text{triplet}} = \max(0, \text{sim}(u, v^-) - \text{sim}(u, v^+) + \text{margin}) \quad (20)$$

Margin: Hyper-parameter for separation



**Figure 4.** SBERT Siamese Network  
[4]

## 2.4. Feature Retrieval

### 2.4.1. Named Entity Recognition

Labelling entities to predefined categories

Input: Barack Obama was born in Hawaii.  
Output: B-PER I-PER O O O B-LOC

- Calculate word embeddings (Hashed Bloom embeddings)
- Apply context (LSTM or iterative CNN)
- Pass through Attention based classifier (Context words and entities as features)

Bloom Embeddings: Multiple hashes

Custom Entities: Calculate embeddings and check for similarity scores across all entities

### 2.4.2. POS Tagging

Identifying the grammatical role of words. Trained similar to NER

Input: The quick brown fox jumps over the lazy dog.  
Output: DT JJ JJ NN VBZ IN DT JJ NN

## 2.5. Other Models

### 2.5.1. Natural Language Inference

Also called Recognizing Textual Entailment (RTE)

Determine relationship between Premise and Hypothesis.

Predict whether hypothesis is Entailed, Contradicted or Neutral.

### 2.5.2. Topic Modelling

Unsupervised approach to identify topics in text

#### Latent Dirichlet Analysis (LDA)

Generative probabilistic model

- Specify number of topics
- Randomly assign topic to each word in each document
- Assign topic to document based on word assignments
- Update document topic probabilities sampling from topic contributions
- Iterate until convergence

$$p(z_{d,n} = t | w_{d,n}, d) \propto \frac{N_{t,d} \cdot N_{w,t}}{N_t} \quad (21)$$

where:

- $N_{t,d}$  is the number of words assigned to topic  $t$  in document  $d$ ,
- $N_{w,t}$  is the number of times word  $w_{d,n}$  is assigned to topic  $t$ ,
- $N_t$  is the total number of words assigned to topic  $t$ .

### Latent Semantic Index (LSI)

Identify hidden topic relations with Singular Value Decomposition

- Select number of topics ( $k$ )
- Build Document Term Frequency matrix
- Perform Singular Value Decomposition
- Extract topic and document vectors for comparison

$$A \approx U \Sigma V^T \quad (22)$$

where:

- $U \in \mathbb{R}^{m \times k}$  is the term-topic matrix,
- $\Sigma \in \mathbb{R}^{k \times k}$  is the diagonal matrix of singular values,
- $V^T \in \mathbb{R}^{k \times n}$  is the document-topic matrix.

### BERTopic

Transformer-based approach

- Document embeddings (BERT)
- Dimensionality Reduction (UMAP)
- Clustering (HDBSCAN)
- Topic Representations (Words with highest TF-IDF scores in each cluster)
- Aggregate cluster for each word in document for topics

Supervised Approach

- Add labels as documents to embed
- Cluster with labels as centroids

## 3. Time Series

### 3.1. Analyses

#### 3.1.1. Decomposition

Time Series Decomposition of Components

$$\text{Additive: } Y_t = T_t + C_t + S_t + R_t$$

$$\text{Multiplicative: } Y_t = T_t \cdot C_t \cdot S_t \cdot R_t \quad (23)$$

$$\text{Multiplicative: } \log Y_t = \log T_t + \log C_t + \log S_t + \log R_t$$

#### $T_t$ : Trend

Long-term movement of data. Calculated by smoothing series

- Moving Averages
- Locally Estimated Scatterplot Smoothing (LOESS) - Regression

#### $S_t$ : Seasonality

Regular fixed-term patterns

- De-Trend the series
- Group by Time Period (January, Monday)
- Calculate average of group
- Center around 0 (Subtract overall mean)

#### $C_t$ : Cyclic

Long-term oscillations not fixed in period

- Remove trend and seasonality
- Smooth the residual data
- Fourier Transform to identify dominant peaks
- Reconstruct Cyclic with dominant peaks

#### $R_t$ : Residuals

Noise. Remainder component.

### 3.1.2. Stationarity

Statistics (mean, variance, covariance) don't change over time

- Strict: Probability Distribution constant
- Weak: Statistic measures remain constant (Common)

#### Augmented Dickey-Fuller Test (ADF)

$H_0: \gamma = 0$  [Unit Root]  $H_1: \gamma < 0$

$$\Delta Y_t = \alpha + \beta t + \gamma Y_{t-1} + \sum_{i=1}^p \delta_i \Delta Y_{t-i} + \varepsilon_t \quad (24)$$

### Differencing

Differencing a time series to stationarize

$$\begin{aligned} \Delta Y_t &= Y_t - Y_{t-1} \text{ (or)} \\ \Delta Y_t &= \frac{Y_t - Y_{t-1}}{Y_{t-1}} \text{ (Log difference)} \end{aligned} \quad (25)$$

### 3.1.3. Autocorrelation

Understand correlation structure

#### Auto-correlation Plot (ACF)

Correlation of value  $Y_t$  with all values  $Y_{t-1} \dots Y_{t-k}$

$$\begin{aligned} \rho_k &= \frac{\sum_{t=1}^{n-k} (Y_t - \bar{Y})(Y_{t+k} - \bar{Y})}{\sum_{t=1}^n (Y_t - \bar{Y})^2} \\ CI &= \pm \frac{t_{\alpha/2}}{\sqrt{n}} \end{aligned} \quad (26)$$

ACF confidence interval increases over lags since fewer observations are available as we look further back.

#### Partial Auto-correlation Plot (PACF)

Correlation of value  $Y_t$  with value  $Y_{t-k}$  without intermediary lags

Fits an OLS model

$$\begin{aligned} Y_t &= \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \dots + \beta_{k-1} Y_{t-(k-1)} + \varepsilon_t \\ \phi_k &= \hat{\beta}_k \\ CI &= \pm \frac{t_{\alpha/2}}{\sqrt{n}} \end{aligned} \quad (27)$$

### 3.2. Forecasting

Fitting time series models and forecasting future values

#### 3.2.1. Exponential Smoothing

Time series with Trend and Seasonality

$$\hat{y}_{t+1} = \alpha y_t + (1 - \alpha) \hat{y}_t \quad (28)$$

Double Exponential (Trend)

$$\begin{aligned} \hat{y}_{t+1} &= L_t + T_t \\ L_t &= \alpha y_t + (1 - \alpha)(L_{t-1} + T_{t-1}) \\ T_t &= \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1} \end{aligned} \quad (29)$$

Triple Exponential Holt Winters (Seasonality)

$$\begin{aligned} \hat{y}_{t+1} &= (L_t + T_t) \cdot S_{t+m} \\ L_t &= \alpha \frac{y_t}{S_{t-p}} + (1 - \alpha)(L_{t-1} + T_{t-1}) \\ T_t &= \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1} \\ S_t &= \gamma \frac{y_t}{L_t} + (1 - \gamma)S_{t-p} \end{aligned} \quad (30)$$

Optimization occurs by Gradient Descent

### 3.2.2. ARIMA

Auto-Regressive Integrated Moving Average  
Fits with MLE estimation or Regression

- Use PACF to get Auto-regressive term (p)
- Apply differencing (d) to get stationary data
- Use ACF to get Moving Average term (q)

$$\begin{aligned}\Phi(B)(1-B)^d Y_t &= \Theta(B)\epsilon_t \\ \text{ARIMA (1,1,1): } Y_t &= \phi_1 Y_{t-1} + (Y_t - Y_{t-1}) + \theta_1 \epsilon_{t-1} + \epsilon_t \\ \text{AR: } Y_t &= \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \phi_p Y_{t-p} + \epsilon_t \\ \text{MA: } Y_t &= \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \dots + \theta_q \epsilon_{t-q}\end{aligned}\quad (31)$$

Residuals should resemble white noise

Assess performance across different ARIMA variants (p, d, q) comparing metrics

**Akaike Information Criteria (AIC)**

$$\text{AIC} = -2 \ln(L) + 2k \quad (32)$$

**Bayesian Information Criteria (BIC)**

$$\text{BIC} = -2 \ln(L) + k \ln(n) \quad (33)$$

- $L$ : Maximum likelihood of the model.
- $k$ : Number of estimated parameters in the model.
- $n$ : Total number of observations in the dataset.

Choose model with a lower AIC / BIC score

**SARIMAX:**

Seasonal ARIMA with Exogenous Variables

$$\Phi_p(B)\Phi_p(B^m)(1-B)^d(1-B^m)^D Y_t = \Theta_q(B)\Theta_q(B^m)\epsilon_t + \beta X_t \quad (34)$$

- Seasonality ( $m$ ): Decompose / Domain knowledge
- Get seasonal p, q from periodic spikes in PACF, ACF
- Seasonal difference subtracts from  $t - m$  observation

Exogenous Variables: Dependent Variables (predictors). Required for future time steps for forecast.

**Table 5.** SARIMAX Components

Component	Variable	Identification
Auto-Regressive	p	Partial Autocorrelation (PACF)
Differencing	d	Augmented Dickey Fuller (ADF)
Moving Average	q	Autocorrelation (ACF)
Seasonality	s	Seasonal Decomposition

### 3.2.3. GARCH

Generalized Auto-Regressive for Conditional Heteroskedasticity  
Models Variance of Time Series

- Fit an ARMA model to the time series
- Fit variance of time series on residuals from ARMA

$$\begin{aligned}Y_t &= \mu + \epsilon_t \\ \sigma_t^2 &= \omega + \sum_{i=1}^q \alpha_i \epsilon_{t-i}^2 + \sum_{j=1}^p \beta_j \sigma_{t-j}^2\end{aligned}\quad (35)$$

Forecasts: Expectation of Errors Variance

$$\begin{aligned}\sigma_{T+h}^2 &= \omega + \sum_{i=1}^q \alpha_i \mathbb{E}[\epsilon_{T+h-i}^2] + \sum_{j=1}^p \beta_j \sigma_{T+h-j}^2 \\ \mathbb{E}[\epsilon_{T+h-i}^2] &= \sigma_{T+h-i}^2\end{aligned}\quad (36)$$

**Selection of (p, q)**

- Fit an initial GARCH (1,1)
- Use ACF, PACF of squared residuals

**Evaluation**

Squared returns act as a good proxy when variance is not available. Can be compared visually / with correlation.

**NOTE:** Conventional ML / RNNs can also be used to forecast stepwise, but lack the inbuilt capabilities to capture more temporal components

## References

- [1] *Activation Functions*. [Online]. Available: <https://medium.com/the-modern-scientist/an-overview-of-activation-functions-in-deep-learning-97a85ac00460>.
- [2] *Neural Networks*. [Online]. Available: <https://www.cs.toronto.edu/~lczhang/360/lec/w02/terms.html>.
- [3] *Recurrent Neural Networks*. [Online]. Available: <https://www.linkedin.com/pulse/rnn-lstm-gru-why-do-we-need-them-suvankar-maity-joegc/>.
- [4] *Siamese*. [Online]. Available: [https://www.researchgate.net/figure/SBERT-Siamese-model-architecture\\_fig3\\_378440592](https://www.researchgate.net/figure/SBERT-Siamese-model-architecture_fig3_378440592).
- [5] *Transformer Architecture*. [Online]. Available: <https://quantdare.com/transformers-is-attention-all-we-need-in-finance-part-i/>.
- [6] *VAE*. [Online]. Available: <https://towardsdatascience.com/uncovering-anomalies-with-variational-autoencoders-vae-a-deep-dive-into-the-world-of-1b2bce47e2e9>.
- [7] *Word2Vec*. [Online]. Available: <https://towardsdatascience.com/nlp-101-word2vec-skip-gram-and-cbow-93512ee24314>.







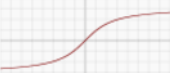


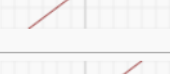
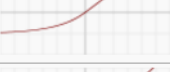
Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$

Figure 5. Activation Functions  
[1]


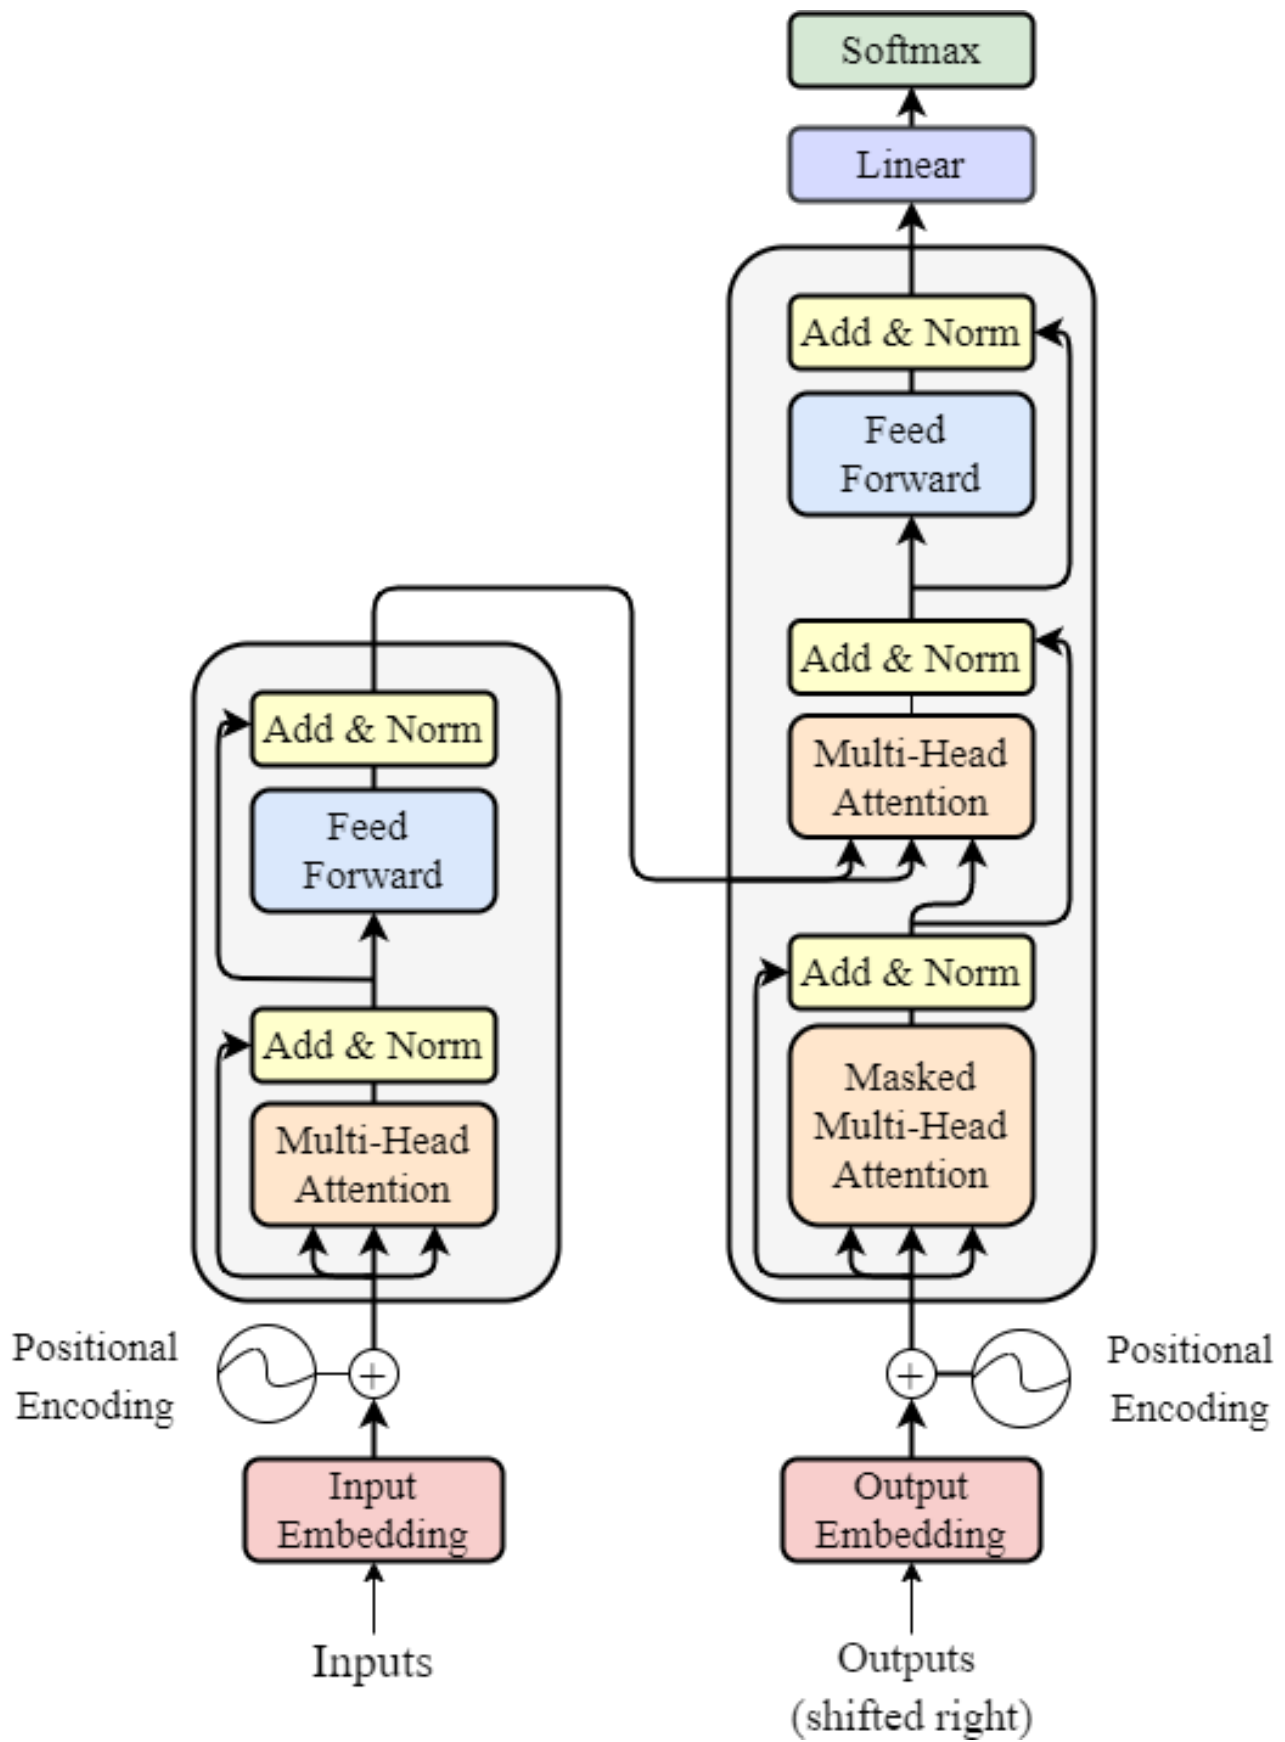


Figure 6. Recurrent Neural Networks  
[3]



**Figure 7.** Attention Based Transformer  
[5]